


Image Representation and Image Processing

CS 3630



Image Processing ...



IN CS, IT CAN BE HARD TO EXPLAIN
THE DIFFERENCE BETWEEN THE EASY
AND THE VIRTUALLY IMPOSSIBLE.

Alt Text:

In the 60s, Marvin Minsky assigned a couple of undergrads to spend the summer programming a computer to use a camera to identify objects in a scene. He figured they'd have the problem solved by the end of the summer. Half a century later, we're still working on it.

Cameras are the primary sensor for most robotic platforms



Common camera types

For now we'll
focus here



Single view
RGB image
(2D data)



Stereo
RGB-D image
(3D data)



Structured light or time of flight
RGB-D image
(3D data)

Images

- An image is basically a 2D array of intensity/color values
- Image types:



Color




Grayscale



B-W


Black and White Images

- A grid (matrix) of 1's and 0's



Grayscale Images

A grid (matrix) of intensity values



=

255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255
255	255	127	145	145	175	127	127	95	47	255	255
255	255	74	127	127	127	95	95	95	47	255	255
255	255	255	74	74	74	74	74	74	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)

Color Images

Three grids (matrices) of intensity values - [R,G,B]




Image Adjustments...

- Scaling
- Color manipulation (color space, color \rightarrow grayscale, etc.)
- Contrast
- Exposure
- ...

Image sampling



640x640



320x320



160x160



80x80



40x40






20x20

Color Models

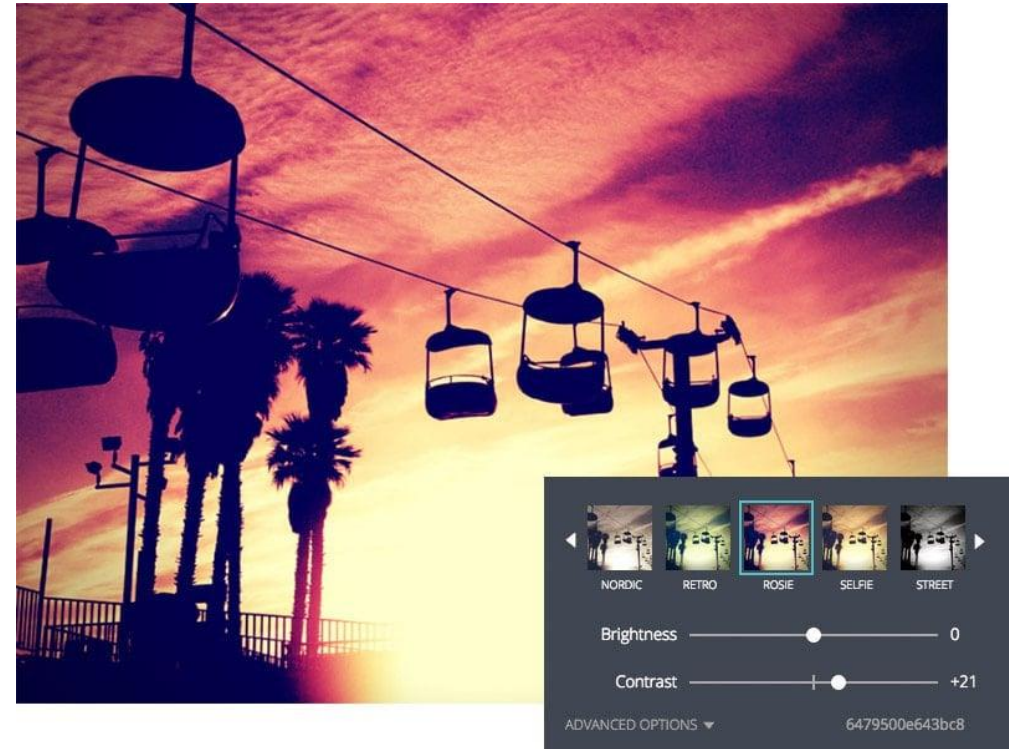
A color model is an abstract mathematical model describing the way colors can be represented as tuples of numbers, typically as three or four values or color components.

- RGB – Red, green, blue
 - CMYK – cyan, magenta, yellow, black
 - HSV – hue, saturation, value
-
- No one color model is always "better" than another.
 - For specific applications, one model might be more suitable than another.
 - e.g., HSV would likely work better when looking for a dark object on top of light background



Filtering


- Types of filtering:
 - Noise removal
 - Edge detection
 - Texture description
 - Multi-scale algorithms
 - Feature detection
 - Matched filters
 - ...



- We will only focus on a few specific methods useful for our application

Monadic operators for filtering

- Operations that take a single pixel as input and output, and do not consider neighboring pixel values



Example: Sepia Tone



$\text{outRed} = (\text{inRed} * .393) + (\text{inGreen} * .769) + (\text{inBlue} * .189)$

$\text{outGreen} = (\text{inRed} * .349) + (\text{inGreen} * .686) + (\text{inBlue} * .168)$

$\text{outBlue} = (\text{inRed} * .272) + (\text{inGreen} * .534) + (\text{inBlue} * .131)$

Example: Histogram Normalization

Input Image




Output Image



What is an image histogram?


For each color C_i , $\mathcal{H}_{C_i}(img)$ provides the number of pixels of color C_i in img .

More generally, a color histogram represents the distribution of various colors (or *intensities* if grayscale) in the image.



Histogram filter

- A histogram can be used to adjust or redistribute the intensity or color distribution of an image
- Usually increases the contrast of an image by effectively spreading out the most frequent intensity values
- Useful in images with backgrounds and foregrounds that are both bright or both dark



Histograms of an image before and after equalization.





Image Contrast

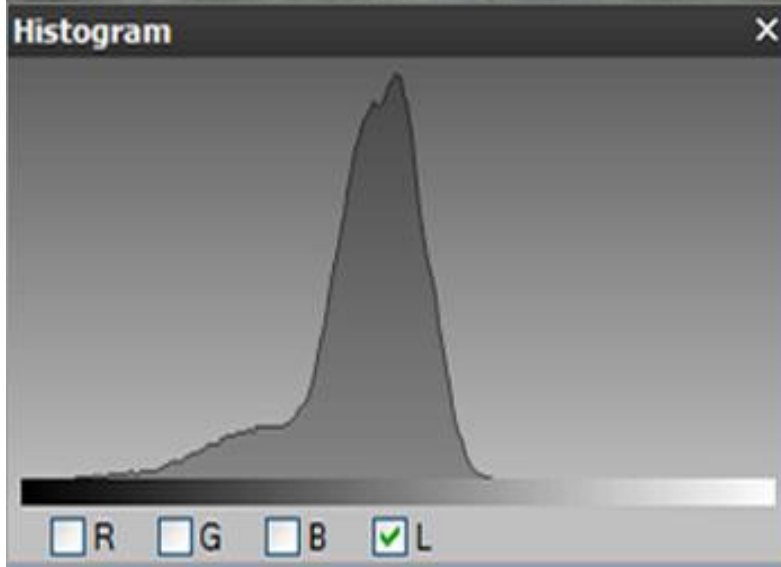
Contrast is the difference between maximum and minimum pixel intensity in a given region



Low contrast


High contrast

Low contrast



Narrow


High contrast




Wide

Example: Histogram Normalization

Input Image




Original Histogram before equalization



Output Image




Equalized histogram



Local operators


Image processing operations that use a region of pixels in the input image to determine the value of a single pixel in the output image



many-to-one mapping between pixels

Example: Linear averaging filter

- Replace each pixel by a linear combination of its neighbors
- The matrix of the linear combination is called the “kernel,” “mask”, or “filter”




What do you think will happen?
Smoothing and blur!

Example: Linear averaging filter


Let F be the image, H be the kernel (of size $2k + 1$ by $2k + 1$), and G be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$




Convolution

- Applying a kernel to an image in this way is called **convolution**.
- NOTE: Convolution can produce unwanted artifacts along the edges of the image.
- Techniques for addressing this include zero padding, edge value replication, mirror extension, and others.



Linear filters: examples



Linear filters: examples



Original



0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Linear filters: examples



Original



0	0	0
1	0	0
0	0	0



Shifted right
By 1 pixel

Linear filters: examples




Original




$\frac{1}{9}$

1	1	1
1	1	1
1	1	1




Mean filter (blurring)

Sharpening



before




after

Sharpening


$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} + \underbrace{\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array}}_{\text{(Original image) - (average of neighbors) = highlights}} = \begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 2 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array}$$

Original image

(Original image) – (average of neighbors) = highlights





before



after

Gaussian Kernel







Approximated by: $\frac{1}{16}$


1	2	1
2	4	2
1	2	1

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$


Gaussian filter



$\sigma = 1$ pixel



$\sigma = 5$ pixels





$\sigma = 10$ pixels





$\sigma = 30$ pixels


Edge Detection




Edge Detection



Origin of Edges



Edges



Edge Detection

Through Convolution

Sobel:

-1	0	1
-2	0	2
-1	0	1



Prewitt:

-1	0	1
-1	0	1
-1	0	1



Canny:

more complex multi-stage algorithm that uses a Gaussian filter and the intensity gradient in an image. One of the most widely-used techniques.



Basic Image Processing

- An image is an array of pixels, which can be binary, grey-value, or color.
- Image filtering takes an image as input and performs basic computation at each pixel to construct an output image.
- Image filters can be used for:
 - Increasing contrast
 - Removing noise
 - Finding edges
- Basic image processing is a precursor for image understanding (aka perception).